

Kadiatou Sidibe

ML OPS 765-01

Homework 1 –Documentation

1. Fork the lab2_factories repo

For this step, I forked the lab2_factories repository to my personal GitHub account to establish my development environment.

In AWS cloud9, I cloned the repository and configured the remote origin to point to my personal fork using my SSH authentication to ensure secure data transfer.

Verification:

I verified the remote configuration to ensure the local environment was correctly linked to my GitHub fork using the following command:

- git remote -v

2. Create an endpoint to dynamically add new topics and store in the topics file

In this step, I implemented a new API endpoint, `POST/topics` to allow registration of new email categories and their descriptions in real-time without requiring manual code changes.

Technical phase:

- I defined a `TopicCreateRequest` schema in `app/api/routes.py` to ensure all incoming new topic includes a mandatory topic name and description.

```
class TopicCreateRequest(BaseModel):  
    name: str  
    description: str
```

- I updated the backend login in `app/api/routes.py` to handle incoming request by reading the existing definitions in `data/topic_keywords.json` file and appending the new topic data.

```

@router.post("/topics")
async def add_topic(request: TopicCreateRequest):
    file_path = "data/topic_keywords.json"

    try:
        with open(file_path, 'r') as f:
            topics = json.load(f)

        topics[request.name] = {"description": request.description}

        with open(file_path, 'w') as f:
            json.dump(topics, f, indent=4)

    return {"status": "success", "message": f"topic '{request.name}' added."}
    except Exception as e:
        raise HTTPException(status_code=500, detail=str(e))

```

Verification:

I initialized the unicorn server and utilized the Swagger UI to execute a “Try it out” test, successfully adding a new topic and verifying in within the JSON database. Detailed results of this verification can be found in the screenshots provided under **Step 5**.

3. Create an endpoint to store emails. These should have an optional ground truth. This ground truth is useful for the similarity classifier

In this step, I implemented a new API endpoint, POST/emails, to store emails.

Technical phase:

- I implemented the `EmailStoreRequest` schema in `app/api/routes.py` to ensure that every stored email includes a subject and body, along with an optional `ground_truth` field.

```

34     class EmailStoreRequest(BaseModel):
35         subject: str
36         body: str
37         ground_truth: str = None
38

```

- I designed the backend logic in `app/api/routes.py` to append these validated entries to the `data/emails.json` file.

```

85     @router.post("/emails")
86     async def store_email(request: EmailStoreRequest):
87         file_path = "data/emails.json"
88
89         try:
90             with open(file_path, 'r') as f:
91                 stored_emails = json.load(f)
92
93             stored_emails.append(request.dict())
94
95             with open(file_path, 'w') as f:
96                 json.dump(stored_emails, f, indent=4)
97
98         return {"status": "success", "message": "Email added to emails.json", "total_emails": len(stored_emails)}
99     except Exception as e:
100         raise HTTPException(status_code=500, detail=f"Error saving email: {str(e)}")
101

```

Verification:

I initialized the uvicorn server and utilized the Swagger UI to execute a “Try it out” test, successfully adding a new email and verifying in the database. Detailed results of this verification can be found in the screenshots provided under **Step 7**.

4. Update the classifier optionally use either the topic classification or to select the class of the most similar email from the stored emails

In this step, I developed an advanced inference mode using Cosine Similarity. When the toggle is enabled, the system generates embeddings for incoming emails and compares them against the stored knowledge base in `data/emails.json`

Technical phase:

- I updated the `app/models/similary_model.py` to include the `predict advanced` which utilizes **Cosine Similarity** to compare input embeddings against the stored knowledge base.

```

def predict_advanced(self, features: Dict[str, Any], use_stored_emails: bool = False) -> str:
    """ Switch between topic classification Or most similar stored email"""
    if not use_stored_emails:
        return self.predict(features)
    # Most similar stored email
    email_embedding = features.get("email_embeddings_average_embedding", None)
    if email_embedding is None or isinstance(email_embedding, list):
        email_embedding = np.array(email_embedding) if email_embedding else None

    stored_emails = self._load_email_data()
    if not stored_emails or email_embedding is None:
        return self.predict(features)

    best_score = -1
    predicted_label = "unknown"

    for example in stored_emails:
        example_embedding = self.model.encode(example['body'], convert_to_numpy=True)

        dot_product = np.dot(email_embedding, example_embedding)
        norm_product = np.linalg.norm(email_embedding) * np.linalg.norm(example_embedding)
        score = dot_product / norm_product if norm_product != 0 else 0

        if score > best_score:
            best_score = score
            predicted_label = example.get('ground_truth', 'unknown')

    return predicted_label

```

- I updated the POST/ emails/classify route to store a bool of use_stored_emails.

```

@router.post("/emails/classify", response_model=EmailClassificationResponse)
async def classify_email(request: EmailRequest, use_stored_emails: bool = False):
    try:
        inference_service = EmailTopicInferenceService()
        email = Email(subject=request.subject, body=request.body)
        result = inference_service.classify_email(email, use_stored_emails = use_stored_emails)
    
```

- I modified app/services/email_topic_inference.py, by updating the classify_email method to pass use_stored_emails boolean through the service layer.

```

def classify_email(self, email: Email, use_stored_emails: bool = False) -> Dict[str, Any]:
    """Classify an email into topics using generated features"""

    # Step 1: Generate features from email
    features = self.feature_factory.generate_all_features(email)

    # Step 2: Classify using features
    predicted_topic = self.model.predict_advanced(features, use_stored_emails)
    topic_scores = self.model.get_topic_scores(features)

    # Return comprehensive results
    return {
        "predicted_topic": predicted_topic,
        "topic_scores": topic_scores,
        "features": features,
        "available_topics": self.model.topics,
        "email": email
    }

```

5. Demonstrate this creating new topics

I conducted the following three steps to confirm the API function work correctly.

The first screenshot below shows a POST request being sent to the /api/v1/topics endpoint. The JSON request name and description.

A screenshot of a POST request to the '/api/v1/topics' endpoint. The request body contains the following JSON:

```
{
  "name": "Refunds",
  "description": "Email regarding customer requests for money back"
}
```

After the execution, the second screenshot verifies that the server processed the request successfully.

- The status code is 200 (Success).
- The response message is “Topic ‘Refunds’ added.”

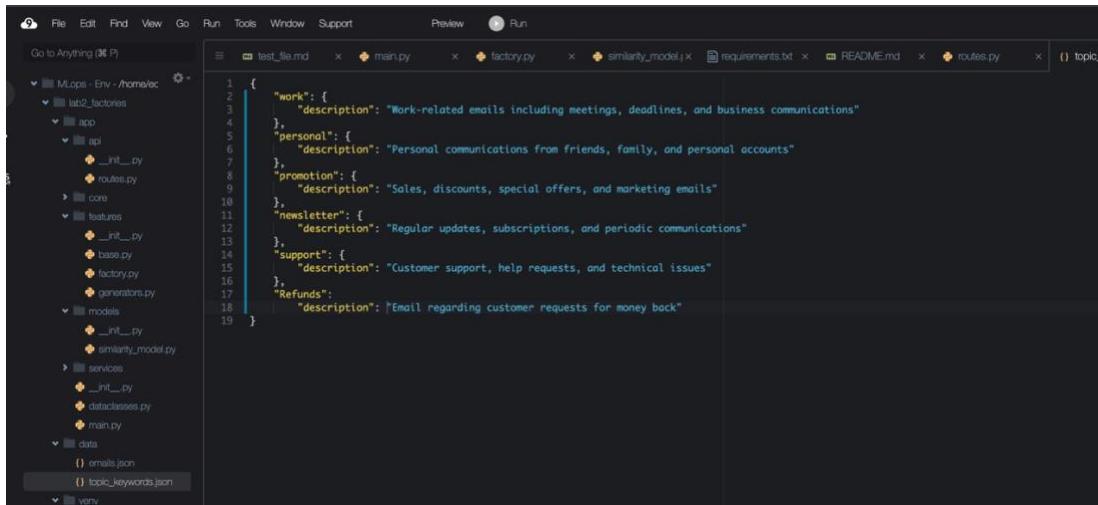
A screenshot of the server response after executing the POST request. The response includes the following details:

- Curl:** curl -X 'POST' \ 'http://54.226.228.68:8000/api/v1/topics' \ -H 'accept: application/json' \ -H 'Content-Type: application/json' \ -d '{ "name": "Refunds", "description": "Email regarding customer requests for money back" }'
- Request URL:** http://54.226.228.68:8000/api/v1/topics
- Server response:**
 - Code:** 200
 - Response body:**

```
{
  "status": "success",
  "message": "topic 'Refunds' added."
}
```
 - Response headers:**

```
content-length: 55
content-type: application/json
date: Sun,22 Feb 2026 02:38:48 GMT
server: uvicorn
```
- Responses:**
 - Code:** 200
 - Description:** Successful Response
 - Media type:** application/json
 - Example Value:** "string"

The third screenshot below shows the content of the `data/topic_keywords.json` file. It confirms that the system successfully stored the refunds topic and description.



```
1  {
2      "work": {
3          "description": "Work-related emails including meetings, deadlines, and business communications"
4      },
5      "personal": {
6          "description": "Personal communications from friends, family, and personal accounts"
7      },
8      "promotion": {
9          "description": "Sales, discounts, special offers, and marketing emails"
10     },
11     "newsletter": {
12         "description": "Regular updates, subscriptions, and periodic communications"
13     },
14     "support": {
15         "description": "Customer support, help requests, and technical issues"
16     },
17     "Refunds": {
18         "description": "Email regarding customer requests for money back"
19     }
}
```

6. Demonstrating this performing inference on the new topics



default

POST /api/v1/emails/classify Classify Email

Parameters

Name	Description
use_stored_emails	false boolean (query)

Request body required

```
{
  "subject": "Where is my stuff",
  "body": "My delivery is missing."
}
```

application/json

Responses

Curl

```
curl -X 'POST' \
  'http://3.94.191.58:8000/api/v1/emails/classify?use_stored_emails=false' \
  -H 'accept: application/json' \
  -H 'Content-Type: application/json' \
  -d '{
    "subject": "Where is my stuff",
    "body": "My delivery is missing."
}'
```

Request URL

http://3.94.191.58:8000/api/v1/emails/classify?use_stored_emails=false

Server response

Code	Details
200	<p>Response body</p> <pre>{ "predicted_topic": "Refunds", "topic_scores": ["work": 0.5679802136305463, "personal": 0.5489057173383705, "promotion": 0.5963446295820656, "newsletter": 0.5706176187353775, "support": 0.626086245693444, "Refunds": 0.618546181722687], "features": { "spam_ham_spam_words": 0, "word_embedding_average_length": 4.25, "email_embedding_average_embedding": [0.012074811848253276, -0.03037879429757595, 0.06938007473945618, 0.030999712312221527, 0.11830092966556549, 0.00445146812160405, -0.000493820363152027, -0.04595845542907715, 0.022034527733922005, 0.019357275996804237, 0.080327956016268302] }</pre>

[Copy](#) [Download](#)

7. Demonstrate this adding new emails

To confirm the API is functioning correctly and data is being persisted, the following three parts processed was performed.

The screenshot below shows a POST request being sent to the `/api/v1//emails` endpoint. The JSON request body includes the subject, body, and the ground truth.

POST /api/v1/emails Store Email

Parameters

No parameters

Request body **required**

application/json

```
{
  "subject": "Missing Package",
  "body": "My order #14567 was never delivered. Please help!",
  "ground_truth": "support"
}
```

Execute

After the tryout is executed, the second screenshot verifies that the server processed the request successfully.

- The status code is 200 (Success).
- The response message is “Email added to `emails.json`”

- It shows the `total_emails` count updated to 1, confirming the backend logic is working.

```

Responses

Curl
curl -X 'POST' \
  'http://54.226.228.68:8000/api/v1/emails' \
  -H 'accept: application/json' \
  -H 'Content-Type: application/json' \
  -d '{
    "subject": "Missing Package",
    "body": "My order #14567 was never delivered. Please help!",
    "ground_truth": "support"
  }'
Request URL
http://54.226.228.68:8000/api/v1/emails
Server response

Code Details
200 Response body
{
  "status": "success",
  "message": "Email added to emails.json",
  "total_emails": 1
}
Response headers
content-length: 76
content-type: application/json
date: Sun, 22 Feb 2026 03:33:59 GMT
server: unicorn

Responses
Code Description
200 Successful Response
Media type
application/json
Controls Accept header.
Links
No links

```

The third screenshot below shows the content of the `data/emails.json` file. It confirms that the system successfully stored the “Missing Package” email.

```

Mlops - Env - /home/ec2-user
  └── lab2_factories
      ├── app
      │   ├── api
      │   │   ├── __init__.py
      │   │   └── routes.py
      │   ├── core
      │   └── features
      │       ├── __init__.py
      │       ├── base.py
      │       ├── factory.py
      │       └── generators.py
      ├── models
      │   ├── __init__.py
      │   └── similarity_model.py
      ├── services
      │   ├── __init__.py
      │   ├── dataclasses.py
      │   └── main.py
      └── data
          ├── emails.json
          └── topic_keywords.json
  └── venv
      ├── bin
      ├── include
      ├── lib
      └── lib64
  └── share
      └── pyvenv.cfg
  └── LICENSE
  └── pyproject.toml
  └── README.md

1: {
2:   "subject": "Missing Package",
3:   "body": "My order #14567 was never delivered. Please help!",
4:   "ground_truth": "support"
5: }
6:
7:

python3 - "ip-172-31-25-21" ✘ Immediate (Javascript (brow ✘ python3 - "ip-172-31-25-21" ✘ bash - "ip-172-31-25-20.e ✘
INFO:  finished server process [44516]
(venv) vclabs:~/environment/lab2_factories (main) $ unicorn app.main:app --host 0.0.0.0 --port 8000
INFO:  Starting server process [40894]
INFO:  Waiting for application to start up.
INFO:  Application startup complete.
INFO:  Unicorn running on http://0.0.0.0:8000 (Press CTRL+C to quit)
INFO:  75.72.46.201:60124 - "GET /docs HTTP/1.1" 200 OK
INFO:  75.72.46.201:60124 - "GET /openapi.json HTTP/1.1" 200 OK
INFO:  75.72.46.201:60138 - "POST /api/v1/emails HTTP/1.1" 200 OK

```

8. Demonstrate this performing inference from the email data

default

POST /api/v1/emails/classify Classify Email

Parameters

Name Description

use_stored_emails true boolean (query)

Request body required

application/json

```
{ "subject": "Where is my stuff", "body": "My delivery is missing." }
```

Responses

Curl

```
curl -X 'POST' \
'http://3.94.191.58:8000/api/v1/emails/classify?use_stored_emails=true' \
-H 'accept: application/json' \
-H 'Content-Type: application/json' \
-d '{
  "subject": "Where is my stuff",
  "body": "My delivery is missing."
}'
```

Request URL

http://3.94.191.58:8000/api/v1/emails/classify?use_stored_emails=true

Server response

Code Details

200 Response body

```
{ "predicted_topic": "support", "topic_scores": { "work": 0.56798802136305463, "personal": 0.5499057173383705, "promotion": 0.5963446295820656, "newsletter": 0.579617618353775, "support": 0.6024886245693644, "Refunds": 0.618546311722687 }, "features": { "spam_has_spam_words": 0, "word_length_average_word_length": 4.25, "email_embeddings_average_embedding": [ 0.012307418184025287, 0.012307418184025287, 0.069238007473945618, 0.030099572312221527 ] } }
```